

✓ Fine Tune NER Model

To fine-tune a Named Entity Recognition (NER) model to extract key entities (products, prices, and location) from Amharic Telegram messages, we will follow these steps.

Step 1: Set Up Environment with GPU Support

- Use Google Colab or GPU-Enabled Environment Ensure that selected a runtime with GPU in Google Colab:
 - Go to Runtime > Change runtime type > Select GPU.
- Install Necessary Libraries
 - Run the following commands in a code cell to install the required libraries: send markdown

✓ Step 1: Import Libraries

```
!pip install pyarrow datasets sequeval
```

```
Requirement already satisfied: pyarrow in /usr/local/lib/python3.11/dist-packages (18.1.0)
Requirement already satisfied: datasets in /usr/local/lib/python3.11/dist-packages (2.14.4)
Collecting sequeval
  Downloading sequeval-1.2.2.tar.gz (43 kB)
    43.6/43.6 kB 3.1 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from datasets) (2.0.2)
Requirement already satisfied: dill<0.3.8,>=0.3.0 in /usr/local/lib/python3.11/dist-packages (from datasets) (0.3.7)
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (from datasets) (2.2.2)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.11/dist-packages (from datasets) (2.32.3)
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.11/dist-packages (from datasets) (4.67.1)
Requirement already satisfied: xxhash in /usr/local/lib/python3.11/dist-packages (from datasets) (3.5.0)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.11/dist-packages (from datasets) (0.70.15)
Requirement already satisfied: fsspec>=2021.11.1 in /usr/local/lib/python3.11/dist-packages (from fsspec[http]>=2021.11.1)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.11/dist-packages (from datasets) (3.11.15)
Requirement already satisfied: huggingface-hub<1.0.0,>=0.14.0 in /usr/local/lib/python3.11/dist-packages (from datasets)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from datasets) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from datasets) (6.0.2)
Requirement already satisfied: scikit-learn>=0.21.3 in /usr/local/lib/python3.11/dist-packages (from sequeval) (1.6.1)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp->dataset)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.11/dist-packages (from aiohttp->datasets) (1.3.1)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp->datasets) (25.3.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.11/dist-packages (from aiohttp->datasets) (1.4.1)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.11/dist-packages (from aiohttp->datasets) (6.0.5)
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp->datasets) (0.2.0)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp->datasets) (1.13.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0.0,>=0.14.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0.0,>=0.14.0)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.2 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0.0,>=0.14.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests>=2.19.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests>=2.19.0->datasets)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests>=2.19.0->datasets)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests>=2.19.0->datasets)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.21.3->sequeval)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.21.3->sequeval)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.21.3->sequeval)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas->datasets)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas->datasets) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas->datasets) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas)
Building wheels for collected packages: sequeval
  Building wheel for sequeval (setup.py) ... done
  Created wheel for sequeval: filename=sequeval-1.2.2-py3-none-any.whl size=16162 sha256=19392f4b1cc7d471a8b4bbce9abfac4b1
  Stored in directory: /root/.cache/pip/wheels/bc/92/f0/243288f899c2eacdffa8c5f9aede4c71a9bad0ee26a01dc5ead
Successfully built sequeval
Installing collected packages: sequeval
Successfully installed sequeval-1.2.2
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_recall_fscore_support
```

```

from transformers import XLRobertaTokenizerFast
from datasets import Dataset, Features, Sequence, Value
from transformers import TrainingArguments
from transformers import XLRobertaForTokenClassification, AutoModelForTokenClassification, AutoTokenizer, Trainer

```

Step 2: Load the Labeled Dataset from CoNLL File

- Load the CoNLL Dataset
 - we can load our CoNLL formatted data into a DataFrame. Here's how we can do that:

```

from google.colab import drive
drive.mount('/content/drive')

```

Mounted at /content/drive

```

# Function to load CoNLL formatted data into a Pandas DataFrame
def load_conll_to_dataframe(file_path):
    """
    Loads data from a CoNLL formatted file into a pandas DataFrame.

    Args:
        file_path (str): The path to the CoNLL file.

    Returns:
        pd.DataFrame: A DataFrame with 'tokens' and 'labels' columns.
    """
    sentences = []
    labels = []
    with open(file_path, 'r', encoding='utf-8') as f:
        sentence_tokens = []
        sentence_labels = []
        for line in f:
            line = line.strip()
            if line: # Non-empty line contains a token and its label
                token, label_item = line.split()
                sentence_tokens.append(token)
                sentence_labels.append(label_item)
            else: # Empty line indicates the end of a sentence
                # Append the completed sentence and its labels
                sentences.append(sentence_tokens)
                labels.append(sentence_labels)
                # Reset lists for the next sentence
                sentence_tokens = []
                sentence_labels = []
        # Append the last sentence if the file does not end with an empty line
        if sentence_tokens:
            sentences.append(sentence_tokens)
            labels.append(sentence_labels)

    # Create a DataFrame from the extracted sentences and labels
    return pd.DataFrame({'tokens': sentences, 'labels': labels})

# Load the CoNLL dataset from the specified file path
df = load_conll_to_dataframe('/content/drive/MyDrive/kaim 6/week4/labeled_data_from_df.conll')

# Explore the first few rows
df.head()

```

	tokens	labels
0	[saachi, electric, kettle, borosilicate, glass...	[O, O, O, O, O, O, O, O, O, O, O, O, B-PRICE, ...
1	[3pcs, bottle, stopper, በማንኛውም, ጠርመሽ, ጭፍ, የሚገጠ...	[O, O, O, O, O, O, O, O, O, O, O, O, O, O, ...
2	[3pcs, bottle, stopper, በማንኛውም, ጠርመሽ, ጭፍ, የሚገጠ...	[O, O, O, O, O, O, O, O, O, O, O, O, O, O, ...
3	[1, pairs, sneaker, crease, protector, ዋጋ, 400...	[B-PRICE, O, O, O, O, B-PRICE, I-PRICE, I-PRIC...
4	[1, pairs, sneaker, crease, protector, ዋጋ, 400...	[B-PRICE, O, O, O, O, B-PRICE, I-PRICE, I-PRIC...

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11457 entries, 0 to 11456
Data columns (total 2 columns):
#   Column   Non-Null Count  Dtype
---  ---
0    tokens   11457 non-null   object
1    labels   11457 non-null   object
dtypes: object(2)
memory usage: 179.1+ KB

```

✓ Define Unique Labels

```

# Get all unique labels from the 'labels' column
all_labels = [label for sublist in df['labels'] for label in sublist]
# Create a set of unique labels to ensure no duplicates
unique_labels = sorted(list(set(all_labels)))

# Create a dictionary mapping each unique label to a unique integer ID
label2id = {label: i for i, label in enumerate(unique_labels)}

# Create a dictionary mapping each integer ID back to its corresponding label
id2label = {i: label for label, i in label2id.items()}

# Print the mappings and the number of unique labels
print("Label to ID mapping:", label2id)
print("ID to Label mapping:", id2label)
print("Number of unique labels:", len(unique_labels))

```

```

Label to ID mapping: {'B-LOC': 0, 'B-PHONE': 1, 'B-PRICE': 2, 'B-PRODUCT': 3, 'I-LOC': 4, 'I-PHONE': 5, 'I-PRICE': 6, 'I-PRODUCT': 7, 'O': 8}
ID to Label mapping: {0: 'B-LOC', 1: 'B-PHONE', 2: 'B-PRICE', 3: 'B-PRODUCT', 4: 'I-LOC', 5: 'I-PHONE', 6: 'I-PRICE', 7: 'I-PRODUCT', 8: 'O'}
Number of unique labels: 9

```

unique_labels

```

['B-LOC',
 'B-PHONE',
 'B-PRICE',
 'B-PRODUCT',
 'I-LOC',
 'I-PHONE',
 'I-PRICE',
 'I-PRODUCT',
 'O']

```

```

df['labels'] = df['labels'].apply(lambda x: [label2id[label] for label in x])

```

✓ Step 3: Convert DataFrame to Hugging Face Dataset

```

# Define the features for the dataset, including tokens and labels
features = Features({
    'tokens': Sequence(Value('string')), # Sequence of strings for tokens
    'labels': Sequence(Value('int64'))    # Sequence of integers for labels (corresponding to label IDs)
})

# Convert DataFrame to Hugging Face Dataset with specified features
dataset = Dataset.from_pandas(df[['tokens', 'labels']], features=features)

# Print the dataset to verify the conversion
print(dataset)

```

```

Dataset({
  features: ['tokens', 'labels'],
  num_rows: 11457
})

```

✓ Step 4: Tokenization and Label Alignment

- Maintain Consistent Tokenization Across Models: When evaluating multiple models, it's important to ensure that each model uses a tokenization method suited to its specific architecture.
- Tokenizers such as those used by XLM-Roberta, DistilBERT, and mBERT differ, so tokenization must be adjusted accordingly for each model.

✓ Tokenize and align label for XLM-ROBERTA

```
# Use the fast tokenizer
# For XLM-Roberta
tokenizer = XLMRobertaTokenizerFast.from_pretrained(
    "xlm-roberta-base",
    clean_up_tokenization_spaces=True
)
# Define tokenizer function
def tokenize_and_align_labels(examples):
    tokenized_inputs = tokenizer(
        examples['tokens'],
        truncation=True,
        is_split_into_words=True,
        padding="max_length", # Padding to max length
        max_length=128 # Adjust as needed
    )
    labels = []
    for i, label in enumerate(examples['labels']):
        word_ids = tokenized_inputs.word_ids(batch_index=i) # Get word ids for each token
        previous_word_idx = None
        label_ids = []
        for word_idx in word_ids:
            if word_idx is None:
                # Token corresponds to special tokens like [CLS], [SEP], etc.
                label_ids.append(-100)
            elif word_idx != previous_word_idx:
                # The first token of a word
                label_ids.append(label[word_idx])
            else:
                # Subword token, assign -100 so it's ignored during training
                label_ids.append(-100)
            previous_word_idx = word_idx
        labels.append(label_ids)
    tokenized_inputs['labels'] = labels
    return tokenized_inputs
# Tokenize the dataset using xlm_roberta
tokenized_xlm_dataset = dataset.map(tokenize_and_align_labels, batched=True)
```

🔄 /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>).
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.

warnings.warn(
tokenizer_config.json: 100% 25.0/25.0 [00:00<00:00, 2.08kB/s]
sentencepiece.bpe.model: 100% 5.07M/5.07M [00:00<00:00, 10.7MB/s]
tokenizer.json: 100% 9.10M/9.10M [00:01<00:00, 6.04MB/s]
config.json: 100% 615/615 [00:00<00:00, 59.0kB/s]
Map: 100% 11457/11457 [00:07<00:00, 2085.59 examples/s]


✓ Tokenize and align for mbert

```
# For mBERT
tokenizer_mbert = AutoTokenizer.from_pretrained(
    'bert-base-multilingual-cased',
    clean_up_tokenization_spaces=True
)
# Define tokenizer function
def tokenize_and_align_labels_mbert(examples):
    tokenized_inputs = tokenizer_mbert(
        examples['tokens'],
        truncation=True,
```

```

        is_split_into_words=True,
        padding="max_length", # Padding to max length
        max_length=128 # Adjust as needed
    )
    labels = []
    for i, label in enumerate(examples['labels']):
        word_ids = tokenized_inputs.word_ids(batch_index=i) # Get word ids for each token
        previous_word_idx = None
        label_ids = []
        for word_idx in word_ids:
            if word_idx is None:
                # Token corresponds to special tokens like [CLS], [SEP], etc.
                label_ids.append(-100)
            elif word_idx != previous_word_idx:
                # The first token of a word
                label_ids.append(label[word_idx])
            else:
                # Subword token, assign -100 so it's ignored during training
                label_ids.append(-100)
            previous_word_idx = word_idx
        labels.append(label_ids)
    tokenized_inputs['labels'] = labels
    return tokenized_inputs
# Tokenize the dataset using mberta
tokenized_mbert_dataset = dataset.map(tokenize_and_align_labels_mbert, batched=True)

```

 tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 1.42kB/s]

config.json: 100% 625/625 [00:00<00:00, 9.77kB/s]

vocab.txt: 100% 996k/996k [00:00<00:00, 4.51MB/s]

tokenizer.json: 100% 1.96M/1.96M [00:00<00:00, 9.05MB/s]

Map: 100% 11457/11457 [00:06<00:00, 2308.74 examples/s]

✓ Tokenize and align label for DistilBERT

```

# For DistilBERT
tokenizer_distilbert = AutoTokenizer.from_pretrained(
    'distilbert-base-multilingual-cased',
    clean_up_tokenization_spaces=True
)
# Define tokenizer function
def tokenize_and_align_labels_distilbert(examples):
    tokenized_inputs = tokenizer_distilbert(
        examples['tokens'],
        truncation=True,
        is_split_into_words=True,
        padding="max_length", # Padding to max length
        max_length=128 # Adjust as needed
    )
    labels = []
    for i, label in enumerate(examples['labels']):
        word_ids = tokenized_inputs.word_ids(batch_index=i) # Get word ids for each token
        previous_word_idx = None
        label_ids = []
        for word_idx in word_ids:
            if word_idx is None:
                # Token corresponds to special tokens like [CLS], [SEP], etc.
                label_ids.append(-100)
            elif word_idx != previous_word_idx:
                # The first token of a word
                label_ids.append(label[word_idx])
            else:
                # Subword token, assign -100 so it's ignored during training
                label_ids.append(-100)
            previous_word_idx = word_idx
        labels.append(label_ids)
    tokenized_inputs['labels'] = labels
    return tokenized_inputs
tokenized_distilbert_dataset = dataset.map(tokenize_and_align_labels_distilbert, batched=True)

```

tokenizer_config.json: 100%	49.0/49.0 [00:00<00:00, 900B/s]
config.json: 100%	466/466 [00:00<00:00, 21.8kB/s]
vocab.txt: 100%	996k/996k [00:00<00:00, 1.50MB/s]
tokenizer.json: 100%	1.96M/1.96M [00:00<00:00, 3.00MB/s]
Map: 100%	11457/11457 [00:04<00:00, 2429.18 examples/s]

✓ Split the dataset into train and test data of each tokenized dataset

```
# Split into train and validation datasets
train_test_split_xlm = tokenized_xlm_dataset.train_test_split(test_size=0.1) # 90% train, 10% validation
train_test_split_mbert = tokenized_mbert_dataset.train_test_split(test_size=0.1) # 90% train, 10% validation
train_test_split_distilbert = tokenized_distilbert_dataset.train_test_split(test_size=0.1) # 90% train, 10% validation
```

✓ Step 5: Set Up Training Arguments

- Configure common training arguments for all models.

```
# Set up training arguments with adjustments
training_args = TrainingArguments(
    output_dir='./results',
    eval_strategy="epoch", # Evaluates at the end of each epoch
    learning_rate=2e-5,
    per_device_train_batch_size=16, # Batch size for training
    per_device_eval_batch_size=16, # Batch size for evaluation
    num_train_epochs=3,
    weight_decay=0.01, # Strength of weight decay
    max_grad_norm=1.0, # Gradient clipping
    logging_dir='./logs', # Directory for storing logs
    logging_strategy="steps", # Log at regular intervals
    logging_steps=50, # Log every 50 steps
    save_strategy="epoch", # Save model at the end of each epoch
    report_to="none", # Only show logs in the output (no TensorBoard)
    # Use GPU for training
    load_best_model_at_end=True, # Load the best model (based on metric) at the end
    metric_for_best_model="eval_loss", # Metric used to determine the best model
    save_total_limit=1, # Only keep the best model, delete the others
)
```

✓ Step 6: Load and Fine-Tune the pre-trained model


- Use Hugging Face Trainer API

Fine-tune the model using the Trainer API.

- fine-tune each of the following pre-trained models:

- xlm-roberta-base
- DistilBERT
- mBERT

```
# Initialize each of the models
# For XLM-Roberta
model_xlmr = XLMRobertaForTokenClassification.from_pretrained("xlm-roberta-base", num_labels=len(unique_labels)) # Ensure ur
# For DistilBERT
model_distilbert = AutoModelForTokenClassification.from_pretrained('distilbert-base-multilingual-cased', num_labels=len(unic
# For mBERT
model_mbert = AutoModelForTokenClassification.from_pretrained('bert-base-multilingual-cased', num_labels=len(unique_labels))
```



model.safetensors: 100%1.12G/1.12G [00:20<00:00, 52.0MB/s]

Some weights of XLMRobertaForTokenClassification were not initialized from the model checkpoint at xlm-roberta-base and You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

model.safetensors: 100%542M/542M [00:17<00:00, 32.5MB/s]

Some weights of DistilBertForTokenClassification were not initialized from the model checkpoint at distilbert-base-multi You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

model.safetensors: 100%714M/714M [00:36<00:00, 20.2MB/s]

Some weights of BertForTokenClassification were not initialized from the model checkpoint at bert-base-multilingual-cas

- Set Up Trainer for Each Model

Create a custom compute_metrics function using scikit-learn to calculate precision, recall, and F1-score, and provide it to each model's Trainer.

```
from seqeval.metrics import precision_score, recall_score, f1_score, classification_report

# Define the function to compute evaluation metrics using seqeval
def compute_metrics(p):
    predictions, labels = p
    predictions = np.argmax(predictions, axis=2)

    # Convert label IDs back to labels for seqeval
    true_labels = [[id2label[label] for label in label_row if label != -100] for label_row in labels]
    predicted_labels = [[id2label[pred] for pred, true in zip(pred_row, true_row) if true != -100] for pred_row, true_row in

    # Use seqeval metrics
    results = {
        "precision": precision_score(true_labels, predicted_labels),
        "recall": recall_score(true_labels, predicted_labels),
        "f1": f1_score(true_labels, predicted_labels)
    }
    return results
```

```
trainer_xlmr = Trainer(
    model=model_xlmr,
    args=training_args,
    train_dataset=train_test_split_xlm['train'],
    eval_dataset=train_test_split_xlm['test'], # Changed from validation to test based on split
    compute_metrics=compute_metrics, # compute f1-score, precision and recall
)
```

```
trainer_distilbert = Trainer(
    model=model_distilbert,
    args=training_args,
    train_dataset=train_test_split_mbert['train'],
    eval_dataset=train_test_split_mbert['test'], # Changed from validation to test based on split
    compute_metrics=compute_metrics, # compute f1-score, precision and recall
)
```

```
trainer_mbert = Trainer(
    model=model_distilbert,
    args=training_args,
    train_dataset=train_test_split_distilbert['train'],
    eval_dataset=train_test_split_distilbert['test'], # Changed from validation to test based on split
    compute_metrics=compute_metrics, # compute f1-score, precision and recall
)
```

Step 8: Model Training, Evaluation, and Comparison:

- Once all three models have been trained, assess their performance using evaluation metrics like F1-score, Precision, Recall, and Accuracy.
- Perform training individually for each model.

```
# Store results of each model
results = {}

# Fine-tune XLM-Roberta
print('For XLM-Roberta')
results['xlmr'] = trainer_xlmr.train()

# Fine-tune DistilBERT
print('For DistilBERT')
results['distilbert'] = trainer_distilbert.train()

# Fine-tune mBERT
print('For mBERT')
results['mbert'] = trainer_mbert.train()

# To get evaluation results after training, you can use trainer.evaluate()
eval_results_xlmr = trainer_xlmr.evaluate()
eval_results_distilbert = trainer_distilbert.evaluate()
eval_results_mbert = trainer_mbert.evaluate()

print("\nXLM-Roberta Evaluation Results:", eval_results_xlmr)
print("DistilBERT Evaluation Results:", eval_results_distilbert)
print("mBERT Evaluation Results:", eval_results_mbert)
```



For XLM-Roberta

[1935/1935 19:15, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Precision	Recall	F1
1	0.002900	0.002833	0.995106	0.998182	0.996642
2	0.002100	0.001358	0.997640	0.999091	0.998365
3	0.001100	0.001125	0.998184	0.999455	0.998819

For DistilBERT

[1935/1935 09:00, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Precision	Recall	F1
1	0.082400	0.079479	0.960484	0.861928	0.908541
2	0.057800	0.054952	0.953181	0.924702	0.938725
3	0.051800	0.045314	0.971851	0.932549	0.951794

For mBERT

[1935/1935 10:17, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Precision	Recall	F1
1	0.041600	0.037899	0.964058	0.939642	0.951693
2	0.036300	0.028982	0.973085	0.959509	0.966249
3	0.029200	0.025556	0.981150	0.962997	0.971988

[72/72 00:07]

[72/72 00:04]

[72/72 00:04]

XLM-Roberta Evaluation Results: {'eval_loss': 0.0011254084529355168, 'eval_precision': 0.9981841292899946, 'eval_recall': 0.9994551104964, 'eval_f1': 0.9988191104964}

DistilBERT Evaluation Results: {'eval_loss': 0.02417793497443199, 'eval_precision': 0.9818656831104964, 'eval_recall': 0.9325491104964, 'eval_f1': 0.9517941104964}

mBERT Evaluation Results: {'eval_loss': 0.02555558830499649, 'eval_precision': 0.9811495673671199, 'eval_recall': 0.9629971104964, 'eval_f1': 0.9719881104964}

- Utilize the Trainer's `evaluate()` method to obtain evaluation metrics for each model, then compare the results side by side to identify which model performs best for the NER task.

✓ Visualize the evaluation result to better understand

```
# Step 1: Evaluate models and collect results
model_results = {
    'XLM-Roberta': trainer_xlmr.evaluate(),
    'DistilBERT': trainer_distilbert.evaluate(),
    'mBERT': trainer_mbert.evaluate()
}

# Step 2: Extract key evaluation metrics
eval_loss = [model_results[model]['eval_loss'] for model in model_results]
```



```
eval_runtime = [model_results[model]['eval_runtime'] for model in model_results]
samples_per_second = [model_results[model]['eval_samples_per_second'] for model in model_results]
model_names = list(model_results.keys())

# Step 3: Visualization (stacked vertically)
plt.figure(figsize=(6, 12))


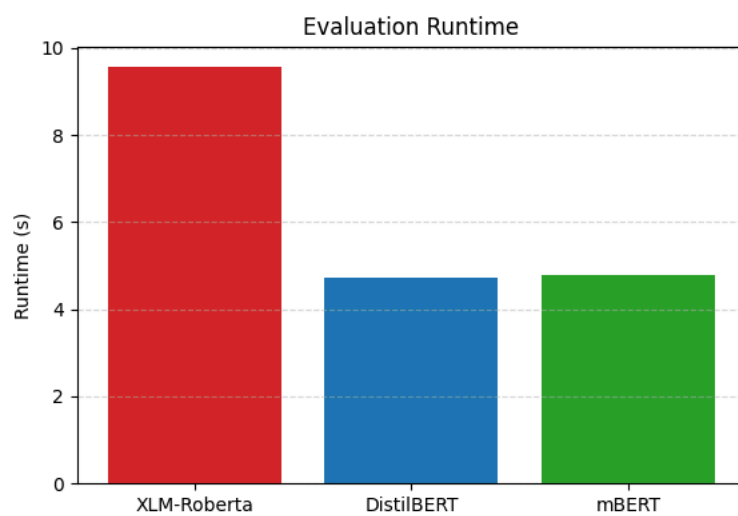

# Evaluation Loss Plot
plt.subplot(3, 1, 1)
plt.bar(model_names, eval_loss, color=['#1f77b4', '#ff7f0e', '#2ca02c'])
plt.title('Evaluation Loss')
plt.ylabel('Loss')
plt.grid(axis='y', linestyle='--', alpha=0.5)

# Evaluation Runtime Plot
plt.subplot(3, 1, 2)
plt.bar(model_names, eval_runtime, color=['#d62728', '#1f77b4', '#2ca02c'])
plt.title('Evaluation Runtime')
plt.ylabel('Runtime (s)')
plt.grid(axis='y', linestyle='--', alpha=0.5)

# Samples Per Second Plot
plt.subplot(3, 1, 3)
plt.bar(model_names, samples_per_second, color=['#1f77b4', '#ff7f0e', '#2ca02c'])
plt.title('Throughput (Samples Per Second)')
plt.ylabel('Samples/sec')
plt.xlabel('Model')
plt.grid(axis='y', linestyle='--', alpha=0.5)

plt.tight_layout()
plt.show()
```

 [72/72 00:09]
 [72/72 05:48]
 [72/72 05:49]



✓ Visualize the result

```

# Step 1: Extract metric values
model_names = list(results.keys())
precision = [results[m]['eval_precision'] for m in model_names]
recall = [results[m]['eval_recall'] for m in model_names]
f1 = [results[m]['eval_f1'] for m in model_names]

# Step 2: Assign distinct colors for each model
colors = {
    'XLM-Roberta': '#1f77b4', # Blue
    'DistilBERT': '#ff7f0e', # Orange
    'mBERT': '#2ca02c' # Green
}
bar_colors = [colors[m] for m in model_names]

# Step 3: Plot vertically
plt.figure(figsize=(7, 12))

# Precision Plot
plt.subplot(3, 1, 1)
  
```

```

bars = plt.bar(model_names, precision, color=bar_colors)
plt.title('Precision by Model')
plt.ylabel('Precision')
plt.ylim(0, 1)
plt.grid(axis='y', linestyle='--', alpha=0.4)
for bar in bars:
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height() + 0.01,
             f'{bar.get_height():.3f}', ha='center', fontsize=9)

# Recall Plot
plt.subplot(3, 1, 2)
bars = plt.bar(model_names, recall, color=bar_colors)
plt.title('Recall by Model')
plt.ylabel('Recall')
plt.ylim(0, 1)
plt.grid(axis='y', linestyle='--', alpha=0.4)
for bar in bars:
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height() + 0.01,
             f'{bar.get_height():.3f}', ha='center', fontsize=9)

# F1-Score Plot
plt.subplot(3, 1, 3)
bars = plt.bar(model_names, f1, color=bar_colors)
plt.title('F1-Score by Model')
plt.ylabel('F1-Score')
plt.ylim(0, 1)
plt.xlabel('Model')
plt.grid(axis='y', linestyle='--', alpha=0.4)
for bar in bars:
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height() + 0.01,
             f'{bar.get_height():.3f}', ha='center', fontsize=9)



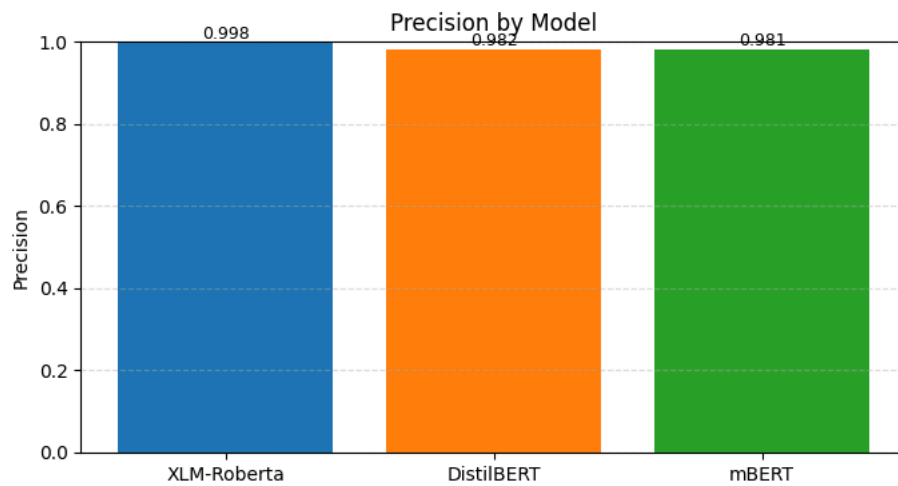
# Add a legend for color meaning
handles = [plt.Rectangle((0, 0), 1, 1, color=colors[m]) for m in model_names]
plt.figlegend(handles, model_names, loc='upper center', ncol=3, frameon=False)

plt.tight_layout(rect=[0, 0, 1, 0.95])
plt.show()

```



■ XLM-Roberta ■ DistilBERT ■ mBERT



```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
from sequeval.metrics import classification_report
import itertools

def plot_confusion_matrix(true_labels, predicted_labels, labels, model_name, ax):
    """Plots a confusion matrix on the given axes."""
    # Flatten labels
    flat_true = list(itertools.chain(*true_labels))
    flat_pred = list(itertools.chain(*predicted_labels))

    cm = confusion_matrix(flat_true, flat_pred, labels=labels)
    cm_df = pd.DataFrame(cm, index=labels, columns=labels)

    sns.heatmap(cm_df, annot=True, cmap='Blues', fmt='g', ax=ax, cbar=False)
    ax.set_title(f'{model_name}')
    ax.set_xlabel('Predicted')
    ax.set_ylabel('True')
```

```

def get_true_and_predicted_labels(trainer, dataset, id2label):
    """Gets predicted and true labels from trainer output."""
    predictions, labels, _ = trainer.predict(dataset)
    predictions = np.argmax(predictions, axis=2)

    true_labels = []
    pred_labels = []

    for label_seq, pred_seq in zip(labels, predictions):
        true_seq = []
        pred_seq_clean = []
        for true_id, pred_id in zip(label_seq, pred_seq):
            if true_id != -100:
                true_seq.append(id2label[true_id])
                pred_seq_clean.append(id2label[pred_id])
        true_labels.append(true_seq)
        pred_labels.append(pred_seq_clean)

    return true_labels, pred_labels

# Get predictions for all models
true_labels_xlmr, pred_labels_xlmr = get_true_and_predicted_labels(trainer_xlmr, train_test_split_xlm['test'], id2label)
true_labels_distilbert, pred_labels_distilbert = get_true_and_predicted_labels(trainer_distilbert, train_test_split_distilbert, id2label)
true_labels_mbert, pred_labels_mbert = get_true_and_predicted_labels(trainer_mbert, train_test_split_mbert['test'], id2label)

# Labels to show (excluding "O")
labels_to_display = sorted([lbl for lbl in id2label.values() if lbl != 'O'])

# Plot all confusion matrices vertically
fig, axes = plt.subplots(3, 1, figsize=(10, 24)) # Adjust height as needed
plot_confusion_matrix(true_labels_xlmr, pred_labels_xlmr, labels_to_display, 'XLM-Roberta', axes[0])
plot_confusion_matrix(true_labels_distilbert, pred_labels_distilbert, labels_to_display, 'DistilBERT', axes[1])
plot_confusion_matrix(true_labels_mbert, pred_labels_mbert, labels_to_display, 'mBERT', axes[2])

plt.tight_layout(rect=[0, 0, 1, 0.98])
plt.suptitle("Confusion Matrices for NER Models", fontsize=18, y=0.995)
plt.show()

# Print detailed classification reports
print("\n📊 XLM-Roberta Classification Report:")
print(classification_report(true_labels_xlmr, pred_labels_xlmr))

print("\n📊 DistilBERT Classification Report:")
print(classification_report(true_labels_distilbert, pred_labels_distilbert))

print("\n📊 mBERT Classification Report:")
print(classification_report(true_labels_mbert, pred_labels_mbert))

```



Confusion Matrices for NER Models

		XLM-Roberta							
True	B-LOC	1202	0	0	0	0	0	0	0
	B-PHONE	0	801	0	0	0	0	0	0
	B-PRICE	0	0	3079	0	0	0	0	0
	B-PRODUCT	0	0	0	418	0	0	0	0
	I-LOC	0	0	0	0	367	0	0	0
	I-PHONE	0	0	0	0	0	356	0	0

✓ Saving the Best model to DRIVE

```
# Identify the best performing model based on f1-score from evaluation results
best_model_name = max(
    ['xlmr', 'distilbert', 'mbert'],
    key=lambda k: {
        'xlmr': eval_results_xlmr,
        'distilbert': eval_results_distilbert,
        'mbert': eval_results_mbert
    }[k]['eval_f1']
)
print(f"The best performing model is: {best_model_name}")

# Save the best model
save_directory = 'best_ner_model'

if best_model_name == 'xlmr':
    trainer_xlmr.save_model(save_directory)
    print(f"XLM-Roberta model saved to {save_directory}")
elif best_model_name == 'distilbert':
    trainer_distilbert.save_model(save_directory)
    print(f"DistilBERT model saved to {save_directory}")
elif best_model_name == 'mbert':
    trainer_mbert.save_model(save_directory)
    print(f"mBERT model saved to {save_directory}")
else:
    print("Could not identify the best model to save.")
```



```
The best performing model is: xlmr
XLM-Roberta model saved to best_ner_model
```

!pip install huggingface-hub